SPEAKER NOTES

Document Processing AI Workshop

📍 SLIDE 1: DOCUMENTS ARE EVERYWHERE

🎯 THE POINT:

Ground them in the reality they already live. Documents aren't going away — they're the lifeblood of operations.

💬 OPEN WITH:

> *"Before we dive into anything technical, I want to start with something you already know — something so obvious it's almost invisible..."*

Build the list:

- Every day, documents flow into your operation — **purchase orders, invoices, bills of lading, delivery receipts**
- On the compliance side — **food safety certifications, audit documentation, quality reports, phytosanitary certificates**
- Here's the stat that matters: **80% of business data doesn't live in your ERP. It lives in documents.**
- And somewhere, right now, someone is **reading a PDF and typing numbers into a system** instead of doing the work they were actually hired for.

❓ ENGAGE THEM:

> *"Think about what flows through YOUR operation every single day. How many different document types? How many hands touch them?"*

➡️ TRANSITION:

> ***"Now — you might be wondering why I'm leading a technology workshop by stating the obvious. Here's why: this 'boring' problem turns out to be the PERFECT fit for AI. Let me show you why..."***

📍 SLIDE 2: THE #1 AI USE CASE

🎯 THE POINT:

Establish credibility immediately. This isn't experimental — it's the most battle-tested AI application in enterprise.

💬 DELIVER WITH CONVICTION:

> *"There's a lot of AI hype out there. Chatbots, image generators, predictions about the future of work. But when you look at where AI is actually delivering MEASURABLE results — not hype, not pilots, ACTUAL production results..."*

Land the punchline:

> *"Document processing is #1. By a wide margin. **76% of enterprises** put this at the top of their automation priorities."*

Explain why it works:

- The ROI is **crystal clear** — you can literally count hours saved and errors eliminated
- It's **low risk** — AI assists humans, it doesn't replace judgment on important decisions
- It's **universal** — every department, every function has documents
- And critically — **the technology is production-ready.** This isn't research. Companies are running this at scale, today.

➡️ TRANSITION:

> ***"So let's get concrete. What does this problem actually LOOK like in a real operation? I'd bet money this sounds familiar..."***

📍 SLIDE 3: THE DOCUMENT CHAOS

🎯 THE POINT:

Make them FEEL the pain they already know. Paint a picture so vivid they're nodding along.

💬 BUILD THE FRUSTRATION:

> *"Let me paint a picture. Tell me if this sounds familiar..."*

Walk through the chaos:

- Invoices arrive as **PDFs attached to emails, sometimes scanned paper, sometimes faxes** — yes, faxes still exist
- Bills of lading are **buried in someone's inbox**, maybe forwarded, maybe not
- Delivery receipts? They're in a **filing cabinet somewhere**, or on a clipboard in a truck, or photographed on someone's phone
- And then **someone has to sit down, read each one, and type that data into your system.** Every. Single. Time.
- The industry average error rate? **4%.** That means out of every 100 documents, 4 have something wrong. And you won't find out until there's a problem.

🌱 GREENHOUSE HOOK:

> *"Meanwhile, there's a **crop that needs scouting**. A **harvest that needs coordinating**. Production work that's NOT getting done because someone's tied to a desk doing data entry."*

❓ PAUSE AND ASK:

> *"Quick mental exercise — how many hours a week does your team spend on paperwork? Just... sit with that number for a second."*

➡️ TRANSITION:

> ***"That frustration you're feeling? It has a dollar sign attached. Let me show you what this actually costs..."***

📍 SLIDE 4: THE REAL NUMBERS

🎯 THE POINT:

Turn frustration into dollars. Let them do the math for their own operation and realize the scale.

📊 DROP THESE NUMBERS:

Metric	Number
Cost per invoice (manual)	**$15-25**
Finance time on data entry	**60%**
Deliveroo before automation	**15 min/invoice**
Deliveroo after	**45 seconds**

💬 MAKE IT REAL:

> *"Let's do some math together. Industry research puts the **fully-loaded cost** of manually processing one invoice at **$15 to $25 dollars.** That's not just salary — that's the back-and-forth, the corrections, the filing, the looking-it-up-later."*

Build to the punchline:

> *"If you're processing 100 invoices a month at $20 each... that's **$24,000 a year.** Just in processing cost. Not the goods — just the PAPERWORK."*

Twist the knife:

- That doesn't count **late payment fees** when things get lost in the shuffle
- It doesn't count the hours spent on **vendor disputes** because someone typed a wrong number
- It's pure **waste.** Labor is expensive, and data entry adds zero value.

➡️ TRANSITION:

> ***"Now — some of you are probably thinking: 'We've tried scanning documents. We looked at OCR. It didn't really work.' Am I right? Let me explain why that happened, and what's different now..."***

📍 SLIDE 5: WHY TRADITIONAL OCR ISN'T ENOUGH

🎯 THE POINT:

Validate their past frustrations. Then show them why this time is different.

💬 ACKNOWLEDGE THEIR SKEPTICISM:

> *"If you've tried document scanning before and been disappointed... you're not alone. Traditional OCR has been around for decades, and for most use cases, it just doesn't cut it."*

Explain the fundamental limitation:

> *"Here's the problem: OCR reads characters. It sees '1-0-0' on a page. But it has NO IDEA if that's a **quantity, a price, a lot number, or part of a date.** It's like hiring someone who can read the letters but **doesn't actually speak the language.**"*

Make it concrete:

- Handwriting? Traditional OCR gets maybe **60% accuracy.** That's almost half wrong.
- Different invoice layouts? It breaks.
- A vendor changes their form? It breaks again.

Pivot to the solution:

> *"Modern AI is fundamentally different. It doesn't just see characters — it UNDERSTANDS context. It knows: 'This is an invoice from Supplier X. The total is HERE. The PO number is THERE. And this handwritten note says...' — it actually comprehends."*

📊 THE PROOF:

Deliveroo achieved 97.6% accuracy with intelligent document processing. Not OCR — intelligent processing.

➡️ TRANSITION:

> ***"So what makes this 'intelligent' processing actually work? Let me break down the architecture..."***

📍 SLIDE 6: INTELLIGENT DOCUMENT PROCESSING

🎯 THE POINT:

IDP is the technical framework. Make it memorable with the "new employee" analogy.

💬 USE THIS ANALOGY:

> *"Think about a new employee on their first day. You hand them an invoice — they can READ it, sure, but they don't know where to put the data, what to check it against, whether it looks right. Now think about that same person after a YEAR. They look at an invoice and instantly know: 'This is from our tomato supplier, the PO matches, the price is normal, route it to Accounts Payable.'"*

Land the insight:

> *"Intelligent Document Processing is like having that **experienced employee** — but one who never takes a day off, never gets tired, and processes documents in seconds instead of minutes."*

Walk through the five steps:

1. **CLASSIFY** — AI identifies: "This is an invoice from Vendor X, not a packing slip"
2. **EXTRACT** — Pulls the relevant data: "Total is $4,500, PO# is 12345, due date is..."
3. **VALIDATE** — Cross-checks: "Does this match what we ordered? Is the price right?"
4. **INTEGRATE** — If everything checks out → auto-posts to your accounting system
5. **LEARN** — Every correction makes it smarter for next time

---

🌱 GREENHOUSE HOOK:

> *"The result? Your team only sees the **exceptions** — the things that actually need human judgment. The routine stuff just... processes."*

➡️ TRANSITION:

> ***"Let me show you what this workflow actually looks like in practice..."***

📍 SLIDE 7: THE IDP WORKFLOW

🎯 THE POINT:

Walk through the 5 steps visually. Use the interactive diagram.

💬 KEY LINES:

> *"You can zoom in on this diagram..."*

Step	What Happens
1. CAPTURE	Document arrives (email, scan, upload)
2. CLASSIFY	AI identifies type, routes accordingly
3. EXTRACT	Pulls relevant fields based on doc type
4. VALIDATE	Cross-checks against your business rules
5. INTEGRATE	Pushes to ERP, triggers workflows

"The magic is in steps 2-4. That's where AI replaces someone sitting at a desk typing."

➡️ TRANSITION:

> ***"What made this actually work well is a recent AI breakthrough..."***

📍 SLIDE 8: THE LLM REVOLUTION

🎯 THE POINT:

This is the "why now" slide. Explain what changed in the last 2 years that made all this actually work.

💬 SET UP THE BREAKTHROUGH:

> *"I need to explain why this is working NOW when similar promises have been made before. Something fundamental changed in 2023."*

Contrast the old way:

> *"Traditional document AI needed **thousands of training examples.** You'd show the system a hundred invoices from one vendor, and MAYBE it could handle that vendor's format. New vendor? Start over. Different layout? Start over. It was expensive, slow, and fragile."*

Reveal the breakthrough:

> *"Large Language Models changed that equation completely. An LLM **already understands** that 'Invoice Total' and 'Amount Due' and 'Grand Total' all mean the same thing. It understands context the way a human does."*

Make it tangible:

- New supplier with a weird invoice format? **The AI figures it out.**
- Handwritten notes on a form? **It reads them in context.**
- Multiple languages in one document? **No problem.**

📊 PROOF POINT:

Wolt — the food delivery company — processes invoices in **9+ languages across 11 countries.** They've had **30% volume growth** with the **same headcount** in AP.

➡️ TRANSITION:

> ***"So with that foundation — what kinds of documents can this technology actually handle? The answer might surprise you..."***

📍 SLIDE 9: WHAT WE CAN PROCESS

🎯 THE POINT:

Expand their imagination. The scope is bigger than they think.

💬 OPEN IT UP:

> *"So what kinds of documents can this actually handle? The short answer: **if a human can read it, AI can too.**"*

Walk through the list:

- ✓ **Invoices and POs** — any format, any vendor, any layout
- ✓ **Bills of lading and shipping documents** — the stuff that comes with every shipment
- ✓ **Delivery receipts** — even with signatures and handwritten notes
- ✓ **Quality certificates and lab reports** — including the weird formatting from third-party labs
- ✓ **Compliance documentation** — food safety certs, audit forms, inspection records
- ✓ **Even handwritten forms** — yes, really. The AI can read handwriting.

🌱 GREENHOUSE SPECIFIC:

> *"For produce operations specifically: phytosanitary certificates, pesticide application records, harvest logs, pack date documentation, temperature logs from cold chain... all of it."*

➡️ TRANSITION:

> ***"Now, here's something important to understand — AI isn't always 100% confident in what it reads. And that's actually a FEATURE, not a bug. Let me explain..."***

📍 SLIDE 10: CONFIDENCE SCORING

🎯 THE POINT:

This is what makes AI trustworthy. It knows what it doesn't know.

💬 EXPLAIN WHY THIS MATTERS:

> *"Here's what makes modern AI different from a black box: **it tells you how confident it is.** Every field, every extraction, comes with a confidence score."*

Walk through the thresholds:

Confidence	What Happens
95%+	Auto-process — no human needed
70-95%	Flag for quick review — human glances at it
<70%	Route to human expert — AI is saying "I'm not sure"

Drive home the insight:

> *"This is crucial for trust. The AI isn't a black box making decisions you can't see.* ***It shows its work.*** *And critically —* ***it knows when to ask for help.****"*

📊 THE PROOF:

- **Adyen** — the global payments company — achieves **93.4% accuracy** on first pass
- With human review on flagged items? They hit **99%+** overall

---

➡️ TRANSITION:

> ***"Which brings us to the question everyone's really thinking about: how do humans fit into this system?"***

📍 SLIDE 11: HUMAN-IN-THE-LOOP

🎯 THE POINT:

Address the unspoken fear in the room. This isn't about replacing people — it's about upgrading what they do.

💬 ADDRESS THE ELEPHANT:

> *"I want to be really clear about something, because I know there's a question in this room that no one's asking out loud: 'Are you talking about replacing my team?'"*

Answer it directly:

> *"No. Here's what actually happens: When the AI is confident, it processes automatically. When it's NOT confident — when something looks unusual or doesn't match —* ***it routes to a human.*** *Your people are still making the decisions that matter. They're just not doing the repetitive typing anymore."*

Explain the virtuous cycle:

- Human reviews the flagged item
- **One click** to approve, correct, or escalate
- Every correction **teaches the system** — it gets smarter
- Over time, fewer things need review

Land the key insight:

> *"Nobody's job becomes 'type data from one screen to another.' That's waste. Their job becomes 'exercise judgment on the things that actually matter.'"*

📊 THE PROOF:

Thermo Fisher runs 53% of their invoices "touchless" — **400,000+ invoices per year** with zero human touch. The other 47%? Reviewed by humans. The humans focus ONLY on what needs their attention.

🌱 GREENHOUSE HOOK:

> *"Imagine your accounting person doing actual ACCOUNTING. Analysis. Vendor negotiations. Cash flow planning. Instead of copying numbers from PDFs."*

➡️ TRANSITION:

> ***"Let's talk about what this looks like in business terms — the actual ROI companies are seeing..."***

📍 SLIDE 12: THE ADOPTION WAVE

🎯 THE POINT:

Create urgency through FOMO. This isn't experimental anymore — it's mainstream.

📊 THE ADOPTION NUMBERS:

Metric	Number
Enterprises using AI	**78%** (up from 55% last year)
Large enterprises with AI	**87%**
IDP market by 2032	**$17.8 billion**

💬 CREATE THE URGENCY:

> *"I want to put this technology in context. This isn't early-adopter territory anymore. This isn't experimental."*

Land the stats:

> *"**78% of enterprises** are now using AI in some form — that's up from 55% just last year. Among large enterprises? **87%.** The IDP market specifically is projected to hit **$17.8 billion** by 2032."*

Reframe the question:

> *"The question isn't WHETHER this technology works. The proof is everywhere. The question is: **do you want to be ahead of your competitors, or behind them?**"*

❓ ENGAGE THE ROOM:

> *"Quick show of hands — who's already using some form of document automation in their operation? Even basic scanning?" (Take note of responses — acknowledge the range)*

➡️ TRANSITION:

> ***"So we know adoption is happening. But what kind of returns are companies actually seeing? Let me show you the numbers..."***

📍 SLIDE 13: THE ROI REALITY

🎯 THE POINT:

Credibility through specifics. These are public numbers from name-brand companies.

📊 THE BIG NUMBERS:

Company	Result
Coca-Cola	€17M savings, 580K hours recovered
Deliveroo	23x efficiency, 15 min → 45 sec
Thermo Fisher	70% time reduction
Payback	5-12 months

💬 DELIVER WITH WEIGHT:

> *"I want to show you some numbers. And I want to be clear — these aren't projections. These aren't 'potential savings.' These are **published results** from companies you've heard of."*

Walk through each one:

- **Coca-Cola's** European operations saved €17 million euros and recovered 580,000 hours of labor
- **Deliveroo** went from 15 minutes per invoice to 45 seconds — that's **23x efficiency**

- **Thermo Fisher** — a $40 billion scientific company — cut document processing time by 70%

Drive home the payback:

> *"And the typical payback period? **5 to 12 months.** This isn't a multi-year initiative. This is something you can see results from THIS YEAR."*

🌱 GREENHOUSE HOOK:

> *"Those hours saved aren't just cost savings on a spreadsheet. They're hours that go back to **production.** To **quality control.** To actually **growing things** instead of processing paperwork."*

➡️ TRANSITION:

> ***"Let me rapid-fire through a few more examples to give you a sense of the range..."***

📍 SLIDE 14: REAL SUCCESS STORIES

🎯 THE POINT:

Name-drop credible companies. Build social proof.

💬 RAPID FIRE:

- **Coca-Cola** — €800M orders, 99% faster
- **JPMorgan** — 360K hours of work → seconds
- **Deutsche Bank** — 16 days → 24 hours
- **Landmark** — 90 min → 4 min per PO
- **Adyen** — 23 countries, 7 people in AP
- **Wolt** — 60% automation, 11 countries, same headcount

> ***"These are Fortune 500 companies. The technology works."***

➡️ TRANSITION:

> ***"Let me paint a picture of the ideal end state..."***

📍 SLIDE 15: THE "TOUCHLESS" INVOICE

🎯 THE POINT:

The goal = invoices that process themselves.

💬 THE FLOW:

1. Invoice arrives by email
2. AI extracts vendor, PO#, line items, totals
3. System matches to purchase order
4. Three-way match: PO vs Invoice vs Receipt
5. If matched → **auto-approve, post, schedule payment**
6. **Human never sees it**

> ***"Your team only sees EXCEPTIONS — the 10% that actually have issues."***

📊 STAT:

Thermo Fisher: 53% touchless = **400,000+ invoices/year** zero human touch

---

➡️ TRANSITION:

> ***"Let's be honest about what this investment looks like..."***

📍 SLIDE 16: COST VS BENEFIT

🎯 THE POINT:

Be honest about costs. Frame as investment vs return.

💬 KEY LINES:

COSTS:

- Initial: software, integration, training (~12 weeks)
- Ongoing: subscription, maintenance

RETURNS:

- Labor savings, faster processing, fewer errors
- Audit readiness, compliance confidence
- Team working on **meaningful tasks**

📊 STAT:

Typical payback: 5-12 months

💡 THE FRAME:

> ***"The question isn't whether you can AFFORD to do this. It's whether you can afford NOT to."***

➡️ TRANSITION:

> ***"Let's talk about practical applications for operations like yours..."***

📍 SLIDE 17: SHIPPING & RECEIVING

🎯 THE POINT:

Bring it to THEIR world. This is the use case they'll see in the demo.

💬 CONNECT TO THEIR REALITY:

> *"Let me bring this back to something very concrete for produce operations: shipping and receiving. This is where documents COLLIDE."*

Paint the complexity:

> *"Think about it — for one shipment, you've got FIVE different data sources that should all agree:"*

Source	What It Is
1. **Purchase Order**	What you ORDERED
2. **Bill of Lading**	What the vendor says they SHIPPED
3. **Invoice**	What the vendor is CHARGING you
4. **Delivery Receipt**	What your team SIGNED for
5. **Actual Count**	What was REALLY there when you opened the boxes

Land the insight:

> *"In a manual world, reconciling these five sources is a headache. It happens late. Discrepancies get discovered weeks later when you're trying to pay the invoice. With AI?* ***All five get compared automatically. Instantly. The moment something doesn't match, you know about it.****"*

🌱 GREENHOUSE HOOK:

> *"Produce operations move FAST. Product is perishable. You can't let a discrepancy sit for a week — by then the evidence is literally rotting."*

➡️ TRANSITION:

> ***"And here's what makes this powerful — when the AI finds a problem, it doesn't just flag it. It routes it INTELLIGENTLY..."***

📍 SLIDE 18: INTELLIGENT ESCALATION

🎯 THE POINT:

AI doesn't just find problems — it routes them to the RIGHT person.

💬 KEY LINES:

Severity	Example	Action
Minor	Ordered 100, got 97	Log it, don't interrupt
Significant	Ordered 100, got 50	Alert receiving manager
Wrong product	Tomato seeds → cucumber	Immediate escalation
Safety issue	Contamination	Notify safety + management

> ***"Not just 'there's a problem' — ACTIONABLE information."***

➡️ TRANSITION:

> ***"This fits into a broader data picture..."***

📍 SLIDE 19: THE DATA FLOW

🎯 THE POINT:

Documents are ONE piece of a bigger unified data system.

💬 KEY LINES:

> *"Documents are one input. Unified insights are the output."*

SOURCES: ERP, climate systems, labor tracking, inventory, shipping **OUTPUT:**

- Leadership → **KPI dashboards**
- Operations → **real-time alerts**
- Growers → **predictions**

➡️ TRANSITION:

> ***"Let me show you the technical architecture..."***

📍 SLIDE 20: DATA ARCHITECTURE

🎯 THE POINT:

For technical folks. Use the interactive diagram.

💬 KEY LINES:

> *"This is the detailed view. You can zoom in to explore."*

Sources → ETL → Data Lake → AI Layer → Users

"Documents aren't siloed anymore. They're part of the same infrastructure as everything else."

➡️ TRANSITION:

> ***"Here's what the transformation looks like..."***

📍 SLIDE 21: KPI TRANSFORMATION

🎯 THE POINT:

Before/after. From chaos to clarity.

💬 KEY LINES:

BEFORE	AFTER
Scattered spreadsheets	Unified data
Manual reconciliation	Automated processing
Reports take days	Same-day visibility

> ***"Nobody's job is pulling data into spreadsheets anymore. The data FLOWS."***

➡️ TRANSITION:

> ***"So how do you actually start?"***

📍 SLIDE 22: GETTING STARTED

🎯 THE POINT:

Remove the "this is too big" objection. Show them the practical first step.

💬 MAKE IT APPROACHABLE:

> *"If you're thinking 'this sounds great but it's a huge undertaking' — let me reframe that. You don't have to automate everything at once. That's not how successful implementations work."*

Walk through the path:

Phase	What Happens
1	**Pick ONE high-volume document type** — usually invoices
2	**Pilot for 30-60 days** — small scope, contained risk
3	**Measure what matters** — accuracy, time saved, exceptions rate
4	**Expand** — add more document types as you prove the model
5	**Integrate** — connect to ERP, automate end-to-end

Set expectations:

> *"From kickoff to production value? **3 to 6 months.** That's not a multi-year transformation. That's a project with a real end date and measurable results."*

💡 THE CALL TO ACTION:

> ***"Start with the pain point that's most acute. Get a win. Build credibility. Then expand from there. That's how this actually works."***

➡️ TRANSITION:

> ***"Alright — we've covered a lot of theory. Now let me show you what this actually LOOKS like when it's running..."***

📍 SLIDE 23: SEE IT IN ACTION

🎯 THE POINT:

Energy shift. Theory is done. Now they get to experience it.

💬 BUILD ANTICIPATION:

> *"Alright. We've covered a lot of ground — the problem, the technology, the results. But talking ABOUT something isn't the same as SEEING it work."*

Set up what's coming:

> *"What you're about to see is a working document processing pipeline. We're going to take a document through the entire flow — capture, classification, extraction, validation, and routing. You'll see the AI's confidence scores. You'll see how exceptions get handled. You'll see the data flow from document to dashboard."*

Make it participatory:

> *"And here's the fun part — **you're going to be part of it.** There will be QR codes you can scan with your phone. You'll submit data. You'll see your inputs flow through the system in real-time. You're not just watching a demo — **you're INSIDE the data pipeline.**"*

🎤 ENERGY CUE:

> ***"So — get your phones out. Make sure you're connected to WiFi. And let's see what intelligent document processing actually looks like when it's running."***

💡 PRESENTER NOTE:

Take a breath here. Let the transition land. Then switch to the demo screen.

📋 QUICK STATS CHEAT SHEET

Stat	Number
Business data in docs	**80%**

Manual invoice cost	**$15-25**
Human error rate	**4%**
Enterprises using AI	**78%**
Coca-Cola savings	**€17M**
Deliveroo efficiency	**23x**
Thermo Fisher touchless	**53%**
Deliveroo accuracy	**97.6%**
Typical payback	**5-12 months**

🌱 GREENHOUSE HOOKS (use when needed)

LABOR: *"Finding good labor is hard enough. Data entry is expensive waste."*

MARGINS: *"Ag runs on tight margins. Manual processing adds ZERO value."*

COMPLIANCE: *"Audits are a fact of life. Searchable docs vs. filing cabinets."*

SEASONALITY: *"Volume spikes at harvest. Manual doesn't scale. Automated does."*

PERISHABILITY: *"Can't let a discrepancy sit for a week. Product's rotted by then."*

REMEMBER: These are talking points, not a script. Build off the key ideas. Make it yours.